

AI Marketing Platform — Buildable Blueprint v2.0

Working name: (TBD) · **Author:** Precious / Metusa Property Ltd · **Date:** August 2026

Status: Pre-build blueprint — replaces “Project Plan & MVP Design v1”

1. Product Overview (Sharpened)

1.1 One-liner

An AI marketing operating system for small businesses and solo founders that turns a single business goal into a complete, executable campaign — strategy, audience, funnel, content, and channel plan — in under 5 minutes.

1.2 The problem (be specific — v1 was too vague)

Small businesses don’t fail at marketing because tools are missing. They fail because:

1. **Strategy is the bottleneck, not content.** ChatGPT can write a post; it can’t tell you *which* campaign to run, in what order, for what budget.
2. **Tools are fragmented.** Strategy lives in someone’s head, content in Canva, scheduling in Buffer, analytics in five dashboards. Nothing closes the loop.
3. **No memory or compounding.** Every campaign starts from zero. Nothing learns what worked last time.

1.3 The wedge (what the MVP actually is)

The MVP is **not** an autonomous agent swarm. It is a **Campaign Generator with a feedback memory**:

Goal in → structured campaign plan + ready-to-publish content out → user reports results → system remembers and improves the next campaign.

Autonomy (auto-publishing, self-optimising agents) is Phase 3+, once the generator proves people will pay for the output.

1.4 Target user (pick one for MVP)

- **Primary:** UK solo founders & micro-businesses (1–10 staff) doing their own

marketing. No marketing hire, £0–500/mo ad budget.

- **Secondary (later):** freelance marketers/agencies running campaigns for multiple clients (this unlocks the Pro tier).

1.5 Positioning vs. competitors

Competitor	What they do	Gap you exploit
Jasper / Copy.ai	Content generation	No strategy layer, no campaign structure, no memory
HubSpot / Mailchimp AI	Channel execution	Heavy, expensive, assumes you already have a strategy
Ocoya / Predis	Social scheduling + AI captions	Post-level, not campaign-level thinking
Agencies	Full service	£1,500+/mo minimum — you're the £49 alternative

Positioning statement: "The marketing director you can't afford, for £49/month."


2. Scope Decisions (what v1 got wrong)

v1 said	v2 decision	Why
5 agents at MVP	1 orchestrated pipeline, agent-shaped later	Agents multiply failure modes; a pipeline is debuggable and 80% of the value
FastAPI backend	Flask (or FastAPI if starting clean)	You already run Flask in production on Metalyzi — reuse auth, Stripe, credit-system patterns
Pinecone vector DB	Supabase pgvector	Already planned for BA Copilot; one DB, no extra vendor, free tier
OpenAI API	Claude API	You have deep prompt-engineering experience with Claude; structured JSON output patterns already

		proven in Metalyzi
Redis cache at MVP	Skip until needed	Premature; Supabase + HTTP caching covers MVP load
LangGraph	Skip at MVP	A sequenced set of Claude calls with JSON contracts is simpler and fully controllable; adopt LangGraph only if Phase 3 agents genuinely need graph state
"Distribution Layer" at MVP	Export + copy-to-clipboard + .zip download	OAuth to Meta/LinkedIn/X is weeks of work and app-review pain; defer to Phase 2

Guiding principle: ship the smallest thing someone will pay £49/month for, on the stack you already know.

3. System Architecture (MVP)



Reuse from Metalyzi (direct ports, ~2 weeks saved):

- Auth + email verification + welcome flow
- Stripe subscription + credit system + post-payment redirect

- Admin dashboard skeleton
 - “AI result card” rendering patterns and JSON-contract prompt style
 - Legal/compliance page templates
-

4. The Campaign Pipeline (replaces “Agent System” at MVP)

One user request triggers a **sequenced pipeline of 5 Claude calls**, each with a strict JSON output contract. Each step’s output feeds the next. Store every intermediate output (for debugging, regeneration of a single step, and future fine-tuning of prompts).

Step 1 — Business Understanding

- **Input:** onboarding data (industry, offer, price point, location, brand voice, past campaign memory if any)
- **Output contract:**

```
{ business_summary, customer_pain_points[], differentiators[], constraints[] }
```
- **Model:** Sonnet

Step 2 — Audience & Positioning

- **Input:** Step 1 output + user goal
- **Output contract:**

```
{ segments[]: { name, demographics, psychographics, watering_holes[], objections[] }, primary_segment, positioning_statement }
```
- **Model:** Sonnet

Step 3 — Funnel & Channel Strategy

- **Input:** Steps 1–2 + budget + timeframe
- **Output contract:**

```
{ funnel_stages[]: { stage, objective, channel, content_types[], kpi, weekly_cadence }, budget_split{}, expected_outcomes }
```
- **Model:** Sonnet

Step 4 — Content Generation (fan-out)

- **Input:** Step 3’s content requirements, brand voice

- **Output contract per asset:** { type, channel, hook, body, cta, hashtags[], image_brief, variant_count }
- Generate 2–3 variants per asset. **Model:** Haiku for short copy (cost control), Sonnet for long-form (email sequences, landing page copy).
- MVP asset types: social posts (IG/FB/LinkedIn/X), ad copy (Meta + Google), 3-email nurture sequence, 1 landing-page copy block.

Step 5 — Campaign Assembly & Calendar

- **Input:** everything above
- **Output contract:** { campaign_name, summary, 30_day_calendar[]: { date, channel, asset_id, action }, launch_checklist[], success_metrics[] }
- **Model:** Sonnet

Memory loop (the moat)

After a campaign runs, the user logs results (manual at MVP: reach, clicks, leads, sales — a 60-second form). Then:

1. A Claude call produces a **campaign retrospective:** { what_worked[], what_failed[], hypotheses[], next_campaign_adjustments[] }
2. Retrospective + campaign summary is embedded (pgvector) into the business's **memory store**.
3. Step 1 of every future campaign retrieves the top-k relevant memories → **campaigns visibly improve over time**. This is the retention mechanic and the demo moment for investors.

Cost estimate per full campaign generation: ~£0.15–0.40 in Claude tokens (Sonnet strategy + Haiku fan-out). At £49/mo with ~10 campaigns/mo, gross margin stays >90%.

5. Database Schema (Supabase)

```
-- users handled by Supabase Auth; profile extension:
profiles          (id uuid PK → auth.users, full_name, plan text, credits int,
                  stripe_customer_id, created_at)

businesses        (id uuid PK, user_id FK, name, industry, offer_description,
```

```

price_point, location, brand_voice jsonb, website, created_at)

campaigns (id uuid PK, business_id FK, goal text, budget numeric,
timeframe_days int, status text
CHECK (status IN ('draft','generating','ready','live','complet
pipeline_outputs jsonb, -- all 5 step outputs, keyed by step
calendar jsonb, created_at, completed_at)

content_assets (id uuid PK, campaign_id FK, type text, channel text,
variant int, content jsonb, image_brief text,
status text CHECK (status IN ('generated','edited','approved','e
edited_content jsonb, created_at)

campaign_results(id uuid PK, campaign_id FK, reported_at, reach int,
clicks int, leads int, sales int, revenue numeric,
user_notes text)

memories (id uuid PK, business_id FK, source_campaign_id FK,
kind text CHECK (kind IN ('retrospective','preference','fact')),
content text, embedding vector(1536), created_at)

credit_ledger (id uuid PK, user_id FK, delta int, reason text,
campaign_id FK NULL, created_at) -- port from Metalyzi

```

Row-Level Security: every table scoped to `user_id` via business ownership — same RLS pattern as Metalyzi.

6. API Endpoints (Flask)

```

POST    /api/auth/*                (port from Metalyzi)
POST    /api/business              create/update business profile
GET     /api/business/:id

POST    /api/campaigns            start generation → returns campaign_id, st
GET     /api/campaigns/:id        poll status + fetch outputs
POST    /api/campaigns/:id/regenerate body: { step: 1-5 } regenerate one step o
POST    /api/campaigns/:id/approve
GET     /api/campaigns?business_id=

GET     /api/campaigns/:id/assets
PATCH  /api/assets/:id           save user edits
POST    /api/assets/:id/regenerate one-asset regen (cheap Haiku call)
GET     /api/campaigns/:id/export .zip of assets + calendar CSV + PDF summar

```

POST	/api/campaigns/:id/results	log results → triggers retrospective + mem
GET	/api/business/:id/memories	"What the AI has learned about your busine
POST	/api/stripe/webhook	(port from Metalyzi)

Generation is async: enqueue on POST, frontend polls GET /campaigns/:id (Metalyzi pattern). Steps stream into `pipeline_outputs` as they finish so the UI can reveal progress step-by-step — great perceived speed.

7. MVP Screens (8 total)

1. **Onboarding wizard** (3 steps): business basics → offer & pricing → brand voice quiz (5 questions, generates `brand_voice` jsonb)
 2. **Dashboard:** campaign list, credits remaining, "New Campaign" CTA, memory count ("Your AI has learned 12 things about your business")
 3. **New Campaign:** single form — goal (free text), budget, timeframe. One button.
 4. **Generation view:** 5-step progress reveal (Step 1 done ✓ → Step 2 generating...) — makes the AI's thinking visible, builds trust
 5. **Campaign view:** strategy summary, audience card, funnel diagram, 30-day calendar
 6. **Asset editor:** asset list with variants, inline editing, per-asset regenerate, approve toggle
 7. **Export:** copy buttons per asset + download bundle (.zip: txt files, calendar .csv, campaign summary .pdf)
 8. **Results & retrospective:** log-results form → AI retrospective card → "applied to future campaigns" confirmation
-

8. Build Plan (6 weeks, solo, Claude Code)

Follows your staged pattern: audit → implement → log → verify → summarise per section.

Week 1 — Foundation

- Repo setup, Supabase schema + RLS, port Metalyzi auth + Stripe + credits

- Ticket size: schema (1d), auth port (2d), Stripe port (1d), skeleton UI (1d)

Week 2 — Pipeline core

- Prompt templates for Steps 1–3 as JSON files (BA Copilot pattern), pipeline runner with per-step persistence, retry + JSON-repair handling
- Verification: 5 test businesses through Steps 1–3, manually score outputs

Week 3 — Content engine

- Step 4 fan-out (Haiku), Step 5 assembly, variant generation, cost logging per campaign
- Verification: full pipeline < 90s, cost per run < £0.40

Week 4 — Frontend

- Screens 1–6, generation progress view, asset editor
- Verification: full user journey clickable end-to-end

Week 5 — Export, results, memory

- Export bundle, results form, retrospective call, pgvector memory write + retrieval into Step 1
- Verification: second campaign for same business demonstrably references first campaign's retrospective

Week 6 — Polish & launch

- Landing page, pricing gates, email flows (port), legal pages (port), 10 beta users from your property/founder network
- Launch channels: your existing social audience, Indie Hackers, UK founder communities

9. Monetisation (validated structure)

Tier	Price	Limits	Target user
Starter	£49/mo	1 business, 5 campaigns/mo, 50 assets	Solo founder
Growth	£149/mo	3 businesses, 20 campaigns/mo, unlimited assets, memory insights	Small business / power user

Pro	£299/mo	10 businesses (client workspaces), white-label exports, priority generation	Freelancers & micro-agencies
-----	---------	---	------------------------------

- **Free trial:** 1 full campaign, no card (the generation view sells itself)
- **Credit overage:** buy extra campaign credits (reuse Metalyzi credit system)
- **Anchor metric for pricing page:** "vs. £1,500/mo for a freelance marketer"

Break-even maths: ~£300/mo infra (Render + Supabase + Vercel + tokens) → 7 Starter customers. 50 customers ≈ £3,500 MRR.

10. Roadmap Beyond MVP

Phase 2 (weeks 7–12) — Distribution: Meta + LinkedIn OAuth publishing, Buffer/Zapier integration as the shortcut, scheduled posting from the calendar, image generation for the `image_brief` fields.

Phase 3 (months 4–6) — Analytics & agents: pull real performance data via channel APIs (replaces manual results form), convert pipeline steps into monitored agents with an optimisation loop, A/B test tracking on variants.

Phase 4 (months 6+) — Compounding moat: cross-customer anonymised benchmarks ("businesses like yours see 2.1% CTR" — the Metalyzi Benchmark Database playbook applied to marketing), agency multi-seat, API access.

11. Risks & Mitigations

Risk	Mitigation
Output quality inconsistent	JSON contracts + per-step regeneration + golden-set of 10 test businesses scored weekly
"ChatGPT can do this" objection	Memory loop + structured campaign (not chat) + calendar/export UX — sell the <i>system</i> , not the text
Token costs blow up	Haiku for fan-out, per-campaign cost logging from day 1, hard credit caps
Channel API rejection (Phase 2)	Buffer/Zapier integration as fallback distribution

Solo-founder bandwidth vs. Metalyzi + BA Copilot	6-week timebox; if no paying beta user by week 8, park it — this plan is designed to be cheap to test
Ad-content compliance (financial/health verticals)	Vertical disclaimer library in prompts; exclude regulated verticals at MVP

12. Definition of Done (MVP)

- A new user goes from signup → complete campaign with ≥ 15 assets in < 10 minutes
- Full pipeline run costs $< \text{£}0.40$ and completes in < 90 seconds
- Second campaign for a business visibly incorporates learnings from the first
- Export bundle opens cleanly and is usable without the app
- Stripe subscription + credit gating live
- 10 beta users onboarded, ≥ 3 have generated 2+ campaigns